

Cheat Sheet: Building Supervised Learning Models

Common supervised learning models

Process Name	Brief Description	Code Syntax
ONE VS ONE CLASSIFIER (USING LOGISTIC REGRESSION)	Process: THIS METHOD TRAINS ONE CLASSIFIER FOR EACH PAIR OF CLASSES. Key hyperparameters: <ul style="list-style-type: none">- 'ESTIMATOR': BASE CLASSIFIER (E.G., LOGISTIC REGRESSION) Pros: CAN WORK WELL FOR SMALL DATASETS. Cons: COMPUTATIONALLY EXPENSIVE FOR LARGE DATASETS. Common applications: MULTICLASS CLASSIFICATION PROBLEMS WHERE THE NUMBER OF CLASSES IS RELATIVELY SMALL.	<pre>from sklearn.multiclass import OneVsOneClassifier from sklearn.linear_model import LogisticRegression model = OneVsOneClassifier(LogisticRegression())</pre>
ONE VS ALL CLASSIFIER (USING LOGISTIC REGRESSION)	Process: TRAINS ONE CLASSIFIER PER CLASS, WHERE EACH CLASSIFIER DISTINGUISHES BETWEEN ONE CLASS AND THE REST. Key hyperparameters: <ul style="list-style-type: none">- 'ESTIMATOR': BASE CLASSIFIER (E.G., LOGISTIC REGRESSION)- 'MULTI_CLASS': STRATEGY TO HANDLE MULTICLASS CLASSIFICATION ('ovr') Pros: SIMPLER AND MORE SCALABLE THAN ONE VS ONE. Cons: LESS ACCURATE FOR HIGHLY IMBALANCED CLASSES. Common applications: COMMON IN MULTICLASS CLASSIFICATION PROBLEMS SUCH AS IMAGE CLASSIFICATION.	<pre>from sklearn.multiclass import OneVsRestClassifier from sklearn.linear_model import LogisticRegression model = OneVsRestClassifier(LogisticRegression())</pre> OR <pre>from sklearn.linear_model import LogisticRegression model_ova = LogisticRegression(multi_class='ovr')</pre>
DECISION TREE CLASSIFIER	Process: A TREE-BASED CLASSIFIER THAT SPLITS DATA INTO SMALLER SUBSETS BASED ON FEATURE VALUES. Key hyperparameters: <ul style="list-style-type: none">- 'MAX_DEPTH': MAXIMUM DEPTH OF THE TREE Pros: EASY TO INTERPRET AND VISUALIZE. Cons: PRONE TO OVERFITTING IF NOT PRUNED PROPERLY. Common applications: CLASSIFICATION TASKS, SUCH AS CREDIT RISK ASSESSMENT.	<pre>from sklearn.tree import DecisionTreeClassifier model = DecisionTreeClassifier(max_depth=5)</pre>
DECISION TREE REGRESSOR	Process: SIMILAR TO THE DECISION TREE CLASSIFIER, BUT USED FOR REGRESSION TASKS TO PREDICT CONTINUOUS VALUES. Key hyperparameters: <ul style="list-style-type: none">- 'MAX_DEPTH': MAXIMUM DEPTH OF THE TREE Pros: EASY TO INTERPRET, HANDLES NONLINEAR DATA. Cons: CAN OVERFIT AND PERFORM POORLY ON NOISY DATA. Common applications: REGRESSION TASKS, SUCH AS PREDICTING HOUSING PRICES.	<pre>from sklearn.tree import DecisionTreeRegressor model = DecisionTreeRegressor(max_depth=5)</pre>
LINEAR SVM CLASSIFIER	Process: A LINEAR CLASSIFIER THAT FINDS THE OPTIMAL HYPERPLANE SEPARATING CLASSES WITH A MAXIMUM MARGIN. Key hyperparameters: <ul style="list-style-type: none">- 'C': REGULARIZATION PARAMETER- 'KERNEL': TYPE OF KERNEL FUNCTION ('linear', 'poly', 'rbf', etc.)- 'GAMMA': KERNEL COEFFICIENT (ONLY FOR 'rbf', 'poly', etc.) Pros: EFFECTIVE FOR HIGH-DIMENSIONAL SPACES. Cons: NOT IDEAL FOR NONLINEAR PROBLEMS WITHOUT KERNEL TRICKS. Common applications: TEXT CLASSIFICATION AND IMAGE RECOGNITION.	<pre>from sklearn.svm import SVC model = SVC(kernel='linear', C=1.0)</pre>
K-NEAREST NEIGHBORS CLASSIFIER	Process: CLASSIFIES DATA BASED ON THE MAJORITY CLASS OF ITS NEAREST NEIGHBORS. Key hyperparameters: <ul style="list-style-type: none">- 'N_NEIGHBORS': NUMBER OF NEIGHBORS TO USE- 'WEIGHTS': WEIGHT FUNCTION USED IN PREDICTION ('uniform' or 'distance')- 'ALGORITHM': ALGORITHM USED TO COMPUTE THE NEAREST NEIGHBORS ('auto', 'ball_tree', 'kd_tree', 'brute') Pros: SIMPLE AND EFFECTIVE FOR SMALL DATASETS. Cons: COMPUTATIONALLY EXPENSIVE AS THE DATASET GROWS. Common applications: RECOMMENDATION SYSTEMS, IMAGE RECOGNITION.	<pre>from sklearn.neighbors import KNeighborsClassifier model = KNeighborsClassifier(n_neighbors=5, weights='uniform')</pre>
RANDOM FOREST REGRESSOR	Process: AN ENSEMBLE METHOD USING MULTIPLE DECISION TREES TO IMPROVE ACCURACY AND REDUCE OVERFITTING. Key hyperparameters: <ul style="list-style-type: none">- 'N_ESTIMATORS': NUMBER OF TREES IN THE FOREST- 'MAX_DEPTH': MAXIMUM DEPTH OF EACH TREE Pros: LESS PRONE TO OVERFITTING THAN INDIVIDUAL DECISION TREES. Cons: MODEL COMPLEXITY INCREASES WITH THE NUMBER OF TREES. Common applications: REGRESSION TASKS SUCH AS PREDICTING SALES OR STOCK PRICES.	<pre>from sklearn.ensemble import RandomForestRegressor model = RandomForestRegressor(n_estimators=100, max_depth=5)</pre>
XGBOOST REGRESSOR	Process: A GRADIENT BOOSTING METHOD THAT BUILDS TREES SEQUENTIALLY TO CORRECT ERRORS FROM PREVIOUS TREES. Key hyperparameters: <ul style="list-style-type: none">- 'N_ESTIMATORS': NUMBER OF BOOSTING ROUNDS- 'LEARNING_RATE': STEP SIZE TO IMPROVE ACCURACY- 'MAX_DEPTH': MAXIMUM DEPTH OF EACH TREE Pros: HIGH ACCURACY AND WORKS WELL WITH LARGE DATASETS. Cons: COMPUTATIONALLY INTENSIVE, COMPLEX TO TUNE. Common applications: PREDICTIVE MODELING, ESPECIALLY IN KAGGLE COMPETITIONS.	<pre>import xgboost as xgb model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1, max_depth=5)</pre>

Associated functions used

Method Name	Brief Description	Code Syntax
OneHotEncoder	TRANSFORMS CATEGORICAL FEATURES INTO A ONE-HOT ENCODED MATRIX.	<pre>from sklearn.preprocessing import OneHotEncoder encoder = OneHotEncoder(sparse=False) encoded_data = encoder.fit_transform(categorical_data)</pre>
accuracy_score	COMPUTES THE ACCURACY OF A CLASSIFIER BY COMPARING PREDICTED AND TRUE LABELS.	<pre>from sklearn.metrics import accuracy_score accuracy = accuracy_score(y_true, y_pred)</pre>
LabelEncoder	ENCODS LABELS (TARGET VARIABLE) INTO NUMERIC FORMAT.	<pre>from sklearn.preprocessing import LabelEncoder encoder = LabelEncoder() encoded_labels = encoder.fit_transform(labels)</pre>
PLOT_TREE	PLOTS A DECISION TREE MODEL FOR VISUALIZATION.	<pre>from sklearn.tree import plot_tree plot_tree(model, max_depth=3, filled=True)</pre>
NORMALIZE	SCALES EACH FEATURE TO HAVE ZERO MEAN AND UNIT VARIANCE (STANDARDIZATION).	<pre>from sklearn.preprocessing import normalize normalized_data = normalize(data, norm='l2')</pre>
COMPUTE_SAMPLE_WEIGHT	COMPUTES SAMPLE WEIGHTS FOR IMBALANCED DATASETS.	<pre>from sklearn.utils.class_weight import compute_sample_weight weights = compute_sample_weight(class_weight='balanced', y=y)</pre>
ROC_AUC_SCORE	COMPUTES THE AREA UNDER THE RECEIVER OPERATING CHARACTERISTIC CURVE (AUC-ROC) FOR BINARY CLASSIFICATION MODELS.	<pre>from sklearn.metrics import roc_auc_score auc = roc_auc_score(y_true, y_score)</pre>

Author

JEFF GROSSMAN
ABHISHEK GAGNEJA



Skills Network